

ML_Project_32

October 13, 2024

ML MID TERM PROJECT - TEAM 32

Project Topic 6 (Hiring Decision)

MS2024013 - Mahima

MT2024146 - Shivani Sanjiv Shukla

MT2024901 - Fitriana Prasari Dewi

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

```
[ ]: df= pd.read_csv('recruitment_data.csv')
df.head()
df.columns
```

```
[ ]: Index(['Age', 'Gender', 'EducationLevel', 'ExperienceYears',
          'PreviousCompanies', 'DistanceFromCompany', 'InterviewScore',
          'SkillScore', 'PersonalityScore', 'RecruitmentStrategy',
          'HiringDecision'],
          dtype='object')
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1500 entries, 0 to 1499
Data columns (total 11 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    1500 non-null   int64
1   Gender                                1500 non-null   int64
2   EducationLevel                        1500 non-null   int64
3   ExperienceYears                       1500 non-null   int64
4   PreviousCompanies                    1500 non-null   int64
5   DistanceFromCompany                  1500 non-null   float64
6   InterviewScore                       1500 non-null   int64
7   SkillScore                           1500 non-null   int64
```

```

8 PersonalityScore      1500 non-null   int64
9 RecruitmentStrategy  1500 non-null   int64
10 HiringDecision      1500 non-null   int64
dtypes: float64(1), int64(10)
memory usage: 129.0 KB

```

```
[ ]: df.describe()
```

```
[ ]:
```

	Age	Gender	EducationLevel	ExperienceYears	\
count	1500.000000	1500.000000	1500.000000	1500.000000	
mean	35.148667	0.492000	2.188000	7.694000	
std	9.252728	0.500103	0.862449	4.641414	
min	20.000000	0.000000	1.000000	0.000000	
25%	27.000000	0.000000	2.000000	4.000000	
50%	35.000000	0.000000	2.000000	8.000000	
75%	43.000000	1.000000	3.000000	12.000000	
max	50.000000	1.000000	4.000000	15.000000	

	PreviousCompanies	DistanceFromCompany	InterviewScore	SkillScore	\
count	1500.000000	1500.000000	1500.000000	1500.000000	
mean	3.00200	25.505379	50.564000	51.116000	
std	1.41067	14.567151	28.626215	29.353563	
min	1.00000	1.031376	0.000000	0.000000	
25%	2.00000	12.838851	25.000000	25.750000	
50%	3.00000	25.502239	52.000000	53.000000	
75%	4.00000	37.737996	75.000000	76.000000	
max	5.00000	50.992462	100.000000	100.000000	

	PersonalityScore	RecruitmentStrategy	HiringDecision
count	1500.000000	1500.000000	1500.000000
mean	49.387333	1.893333	0.310000
std	29.353201	0.689642	0.462647
min	0.000000	1.000000	0.000000
25%	23.000000	1.000000	0.000000
50%	49.000000	2.000000	0.000000
75%	76.000000	2.000000	1.000000
max	100.000000	3.000000	1.000000

```
[ ]: df.isnull().sum()
```

```
[ ]: Age          0
Gender          0
EducationLevel  0
ExperienceYears 0
PreviousCompanies 0
DistanceFromCompany 0
InterviewScore  0
SkillScore      0
```

```

PersonalityScore      0
RecruitmentStrategy    0
HiringDecision         0
dtype: int64

```

```

[ ]: X_=df.drop("HiringDecision",axis=1) 0.8966666666666666
Y_=df["HiringDecision"]
df.head()

```

```

[ ]:
  Age  Gender  EducationLevel  ExperienceYears  PreviousCompanies  \
0   26      1              2              0              3
1   39      1              4              12              3
2   48      0              2              3              2
3   34      1              2              5              2
4   30      0              1              6              1

  DistanceFromCompany  InterviewScore  SkillScore  PersonalityScore  \
0          26.783828          48          78          91
1          25.862694          35          68          80
2           9.920805          20          67          13
3           6.407751          36          27          70
4          43.105343          23          52          85

  RecruitmentStrategy  HiringDecision
0                   1              1
1                   2              1
2                   2              0
3                   3              0
4                   2              0

```

```

[ ]: X_.head()

```

```

[ ]:
  Age  Gender  EducationLevel  ExperienceYears  PreviousCompanies  \
0   26      1              2              0              3
1   39      1              4              12              3
2   48      0              2              3              2
3   34      1              2              5              2
4   30      0              1              6              1

  DistanceFromCompany  InterviewScore  SkillScore  PersonalityScore  \
0          26.783828          48          78          91
1          25.862694          35          68          80
2           9.920805          20          67          13
3           6.407751          36          27          70
4          43.105343          23          52          85

  RecruitmentStrategy


```

0	1
1	2
2	2
3	3
4	2

```
[ ]: Y_.head()
```

```
[ ]: 0    1  
     1    1  
     2    0  
     3    0  
     4    0  
     Name: HiringDecision, dtype: int64
```

```
[ ]: import seaborn as sns  
     plt.figure(figsize=(12,10))  
     cor=X_.corr()  
     sns.heatmap(cor,annot=True,cmap=plt.cm.CMRmap_r)  
     plt.show()
```



```
[ ]: def correlation(dataset, threshold):
    col_corr = set() # Set of all the names of correlated columns
    corr_matrix = dataset.corr()
    for i in range(len(corr_matrix.columns)):
        for j in range(i):
            if abs(corr_matrix.iloc[i, j]) > threshold: # we are interested in
↳absolute coeff value
                colname = corr_matrix.columns[i] # getting the name of column
                col_corr.add(colname)
    return col_corr
corr_feature=correlation(X_,0.7)
len(set(corr_feature))
```

```
[ ]: 0
```

```
[ ]: def standardize_column(column):
    mean = np.mean(column)
    std_dev = np.std(column)
    return (column - mean) / std_dev
```

```
[ ]: df_standardized = df.copy() # Make a copy of the original DataFrame
for column in df.columns:
    df_standardized[column] = standardize_column(df[column])
print("Original Data:")
print(df)
print("\nStandardized Data (Z-score normalization):")
print(df_standardized)
```

Original Data:

	Age	Gender	EducationLevel	ExperienceYears	PreviousCompanies	\
0	26	1	2	0	3	
1	39	1	4	12	3	
2	48	0	2	3	2	
3	34	1	2	5	2	
4	30	0	1	6	1	
...	
1495	48	0	2	3	4	
1496	27	1	2	10	3	
1497	24	1	1	1	2	
1498	48	0	2	4	4	
1499	34	1	2	11	5	

	DistanceFromCompany	InterviewScore	SkillScore	PersonalityScore	\
0	26.783828	48	78	91	
1	25.862694	35	68	80	
2	9.920805	20	67	13	
3	6.407751	36	27	70	
4	43.105343	23	52	85	
...	
1495	9.183783	66	3	80	
1496	14.847731	43	97	7	
1497	4.289911	31	91	58	
1498	36.299263	9	37	44	
1499	12.910472	63	40	26	

	RecruitmentStrategy	HiringDecision
0	1	1
1	2	1
2	2	0
3	3	0
4	2	0
...
1495	3	1

1496	2	0
1497	1	1
1498	2	1
1499	2	1

[1500 rows x 11 columns]

Standardized Data (Z-score normalization):

	Age	Gender	EducationLevel	ExperienceYears	PreviousCompanies	\
0	-0.989083	1.016130	-0.218057	-1.658237	-0.001418	
1	0.416376	1.016130	2.101694	0.928044	-0.001418	
2	1.389387	-0.984126	-0.218057	-1.011667	-0.710538	
3	-0.124185	1.016130	-0.218057	-0.580620	-0.710538	
4	-0.556634	-0.984126	-1.377932	-0.365097	-1.419657	
...	
1495	1.389387	-0.984126	-0.218057	-1.011667	0.707701	
1496	-0.880971	1.016130	-0.218057	0.496997	-0.001418	
1497	-1.205308	1.016130	-1.377932	-1.442714	-0.710538	
1498	1.389387	-0.984126	-0.218057	-0.796144	0.707701	
1499	-0.124185	1.016130	-0.218057	0.712520	1.416821	

	DistanceFromCompany	InterviewScore	SkillScore	PersonalityScore	\
0	0.087792	-0.089598	0.916174	1.418126	
1	0.024537	-0.543879	0.575386	1.043255	
2	-1.070200	-1.068049	0.541307	-1.240051	
3	-1.311444	-0.508934	-0.821844	0.702463	
4	1.208598	-0.963215	0.030126	1.213651	
...	
1495	-1.120812	0.539406	-1.639734	1.043255	
1496	-0.731866	-0.264321	1.563671	-1.444526	
1497	-1.456877	-0.683657	1.359198	0.293513	
1498	0.741221	-1.452440	-0.481056	-0.183596	
1499	-0.864899	0.434572	-0.378820	-0.797022	

	RecruitmentStrategy	HiringDecision
0	-1.295790	1.491914
1	0.154721	1.491914
2	0.154721	-0.670280
3	1.605233	-0.670280
4	0.154721	-0.670280
...
1495	1.605233	1.491914
1496	0.154721	-0.670280
1497	-1.295790	1.491914
1498	0.154721	1.491914
1499	0.154721	1.491914

[1500 rows x 11 columns]

```
[ ]: def min_max_normalize_column(column):
    min_val = np.min(column)
    max_val = np.max(column)
    return (column - min_val) / (max_val - min_val)

# Apply Min-Max normalization to each feature column
df_min_max = df.copy() # Make a copy of the original DataFrame
for column in df.columns:
    df_min_max[column] = min_max_normalize_column(df[column])

print("\nMin-Max Normalized Data (0-1 range):")
print(df_min_max)
```

Min-Max Normalized Data (0-1 range):

	Age	Gender	EducationLevel	ExperienceYears	PreviousCompanies	\
0	0.200000	1.0	0.333333	0.000000		0.50
1	0.633333	1.0	1.000000	0.800000		0.50
2	0.933333	0.0	0.333333	0.200000		0.25
3	0.466667	1.0	0.333333	0.333333		0.25
4	0.333333	0.0	0.000000	0.400000		0.00
...	
1495	0.933333	0.0	0.333333	0.200000		0.75
1496	0.233333	1.0	0.333333	0.666667		0.50
1497	0.133333	1.0	0.000000	0.066667		0.25
1498	0.933333	0.0	0.333333	0.266667		0.75
1499	0.466667	1.0	0.333333	0.733333		1.00

	DistanceFromCompany	InterviewScore	SkillScore	PersonalityScore	\
0	0.515450	0.48	0.78	0.91	
1	0.497013	0.35	0.68	0.80	
2	0.177927	0.20	0.67	0.13	
3	0.107611	0.36	0.27	0.70	
4	0.842135	0.23	0.52	0.85	
...	
1495	0.163175	0.66	0.03	0.80	
1496	0.276542	0.43	0.97	0.07	
1497	0.065221	0.31	0.91	0.58	
1498	0.705907	0.09	0.37	0.44	
1499	0.237767	0.63	0.40	0.26	

	RecruitmentStrategy	HiringDecision
0	0.0	1.0
1	0.5	1.0
2	0.5	0.0
3	1.0	0.0
4	0.5	0.0
...

1495	1.0	1.0
1496	0.5	0.0
1497	0.0	1.0
1498	0.5	1.0
1499	0.5	1.0

[1500 rows x 11 columns]

```
[ ]: X = df_min_max.drop("HiringDecision",axis=1)
Y = df_min_max["HiringDecision"]
df_min_max.head()
```

```
[ ]:      Age  Gender  EducationLevel  ExperienceYears  PreviousCompanies  \
0  0.200000    1.0         0.333333         0.000000             0.50
1  0.633333    1.0         1.000000         0.800000             0.50
2  0.933333    0.0         0.333333         0.200000             0.25
3  0.466667    1.0         0.333333         0.333333             0.25
4  0.333333    0.0         0.000000         0.400000             0.00

      DistanceFromCompany  InterviewScore  SkillScore  PersonalityScore  \
0              0.515450             0.48         0.78              0.91
1              0.497013             0.35         0.68              0.80
2              0.177927             0.20         0.67              0.13
3              0.107611             0.36         0.27              0.70
4              0.842135             0.23         0.52              0.85

      RecruitmentStrategy  HiringDecision
0              0.0         1.0
1              0.5         1.0
2              0.5         0.0
3              1.0         0.0
4              0.5         0.0
```

0.1 Train Test Split

```
[ ]: from sklearn.model_selection import train_test_split
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.20)
```

```
[ ]: print(X_train.shape)
print(Y_train.shape)
```

```
(1200, 10)
(1200,)
```

0.2 Logistic Regression

```
[ ]: import numpy as np

# Sigmoid function (maps values between 0 and 1)
def sigmoid(z):
    return 1 / (1 + np.exp(-z))

# Logistic Regression model from scratch
class LogisticRegressionScratch:
    def __init__(self, learning_rate=0.01, iterations=10000):
        self.learning_rate = learning_rate
        self.iterations = iterations
        self.weights = None
        self.bias = None

    def fit(self, X, y):
        # Number of samples and features
        n_samples, n_features = X.shape

        # Initialize weights and bias
        self.weights = np.zeros(n_features)
        self.bias = 0

        # Gradient descent
        for _ in range(self.iterations):
            # Linear model ( $z = wX + b$ )
            linear_model = np.dot(X, self.weights) + self.bias
            # Apply sigmoid function to get predictions
            y_predicted = sigmoid(linear_model)

            # Compute gradients for weights and bias
            dw = (1 / n_samples) * np.dot(X.T, (y_predicted - y))
            db = (1 / n_samples) * np.sum(y_predicted - y)

            # Update weights and bias using gradient descent
            self.weights -= self.learning_rate * dw
            self.bias -= self.learning_rate * db

    def predict(self, X):
        # Linear model ( $z = wX + b$ )
        linear_model = np.dot(X, self.weights) + self.bias
        # Apply sigmoid function
        y_predicted = sigmoid(linear_model)
        # Convert probabilities to binary classes (0 or 1)
        y_predicted_class = [1 if i > 0.5 else 0 for i in y_predicted]
        return np.array(y_predicted_class)
```

```

# Example usage
if __name__ == "__main__":
    # Create a dataset (replace with your actual data)
    x_train = X_train
    y_train = Y_train

    # Initialize the logistic regression model
    model = LogisticRegressionScratch(learning_rate=0.01, iterations=10000)
    # Train the model
    model.fit(x_train, y_train)
    # Make predictions
    predictions = model.predict(x_train)

    print("Predicted classes:", predictions)

```

Predicted classes: [1 1 0 ... 0 0 1]

0.3 Decision trees

```

[ ]: class Node():
    def __init__(self, feature_index=None, threshold=None, left=None,
    ↪right=None, info_gain=None, value=None):
        ''' constructor '''

        # for decision node
        self.feature_index = feature_index
        self.threshold = threshold
        self.left = left
        self.right = right
        self.info_gain = info_gain

        # for leaf node
        self.value = value

```

```

[ ]: class DecisionTreeClassifier():
    def __init__(self, min_samples_split=2, max_depth=4):
        ''' constructor '''

        # initialize the root of the tree
        self.root = None

        # stopping conditions
        self.min_samples_split = min_samples_split
        self.max_depth = max_depth

    def build_tree(self, dataset, curr_depth=0):

```

```

''' recursive function to build the tree '''

X, Y = dataset[:, :-1], dataset[:, -1]
num_samples, num_features = np.shape(X)

# split until stopping conditions are met
if num_samples >= self.min_samples_split and curr_depth <= self.max_depth:
    # find the best split
    best_split = self.get_best_split(dataset, num_samples, num_features)
    # check if information gain is positive
    if best_split["info_gain"] > 0:
        # recur left
        left_subtree = self.build_tree(best_split["dataset_left"],
    ↪ curr_depth+1)
        # recur right
        right_subtree = self.build_tree(best_split["dataset_right"],
    ↪ curr_depth+1)
        # return decision node
        return Node(best_split["feature_index"],
    ↪ best_split["threshold"],
                        left_subtree, right_subtree,
    ↪ best_split["info_gain"])

    # compute leaf node
    leaf_value = self.calculate_leaf_value(Y)
    # return leaf node
    return Node(value=leaf_value)
def get_best_split(self, dataset, num_samples, num_features):
    ''' function to find the best split '''

    # dictionary to store the best split
    best_split = {}
    max_info_gain = -float("inf")

    # loop over all the features
    for feature_index in range(num_features):
        feature_values = dataset[:, feature_index]
        possible_thresholds = np.unique(feature_values)
        # loop over all the feature values present in the data
        for threshold in possible_thresholds:
            # get current split
            dataset_left, dataset_right = self.split(dataset,
    ↪ feature_index, threshold)
            # check if childs are not null
            if len(dataset_left) > 0 and len(dataset_right) > 0:
                y, left_y, right_y = dataset[:, -1], dataset_left[:, -1],
    ↪ dataset_right[:, -1]

```

```

        # compute information gain
        curr_info_gain = self.information_gain(y, left_y, right_y,
↪ "gini")

        # update the best split if needed
        if curr_info_gain > max_info_gain:
            best_split["feature_index"] = feature_index
            best_split["threshold"] = threshold
            best_split["dataset_left"] = dataset_left
            best_split["dataset_right"] = dataset_right
            best_split["info_gain"] = curr_info_gain
            max_info_gain = curr_info_gain

    # return best split
    return best_split

def split(self, dataset, feature_index, threshold):
    ''' function to split the data '''

    dataset_left = np.array([row for row in dataset if
↪ row[feature_index] <= threshold])
    dataset_right = np.array([row for row in dataset if
↪ row[feature_index] > threshold])
    return dataset_left, dataset_right

def information_gain(self, parent, l_child, r_child, mode="entropy"):
    ''' function to compute information gain '''

    weight_l = len(l_child) / len(parent)
    weight_r = len(r_child) / len(parent)
    if mode == "gini":
        gain = self.gini_index(parent) - (weight_l * self.
↪ gini_index(l_child) + weight_r * self.gini_index(r_child))
    else:
        gain = self.entropy(parent) - (weight_l * self.entropy(l_child) +
↪ weight_r * self.entropy(r_child))
    return gain

def entropy(self, y):
    ''' function to compute entropy '''

    class_labels = np.unique(y)
    entropy = 0
    for cls in class_labels:
        p_cls = len(y[y == cls]) / len(y)
        entropy += -p_cls * np.log2(p_cls)
    return entropy

```

```

def gini_index(self, y):
    ''' function to compute gini index '''

    class_labels = np.unique(y)
    gini = 0
    for cls in class_labels:
        p_cls = len(y[y == cls]) / len(y)
        gini += p_cls**2
    return 1 - gini

def calculate_leaf_value(self, Y):
    ''' function to compute leaf node '''

    Y = list(Y)
    return max(Y, key=Y.count)

def print_tree(self, tree=None, indent=" "):
    ''' function to print the tree '''

    if not tree:
        tree = self.root

    if tree.value is not None:
        print(tree.value)

    else:
        print("X_"+str(tree.feature_index), "<=", tree.threshold, "?", tree.
↳info_gain)
        print("%sleft:" % (indent), end="")
        self.print_tree(tree.left, indent + indent)
        print("%sright:" % (indent), end="")
        self.print_tree(tree.right, indent + indent)

def fit(self, X, Y):
    ''' function to train the tree '''

    dataset = np.concatenate((X, Y), axis=1)
    self.root = self.build_tree(dataset)

def predict(self, X):
    ''' function to predict new dataset '''

    predictions = [self.make_prediction(x, self.root) for x in X]
    return predictions

def make_prediction(self, x, tree):
    ''' function to predict a single data point '''

```

```

        if tree.value!=None: return tree.value
        feature_val = x[tree.feature_index]
        if feature_val<=tree.threshold:
            return self.make_prediction(x, tree.left)
        else:
            return self.make_prediction(x, tree.right)

```

0.4 Training on Both models

```

[ ]: predictions = model.predict(X_test)
     print(len(predictions))

```

300

```

[ ]: from sklearn.metrics import accuracy_score, confusion_matrix, \
     ↪classification_report
     accuracy = accuracy_score(Y_test, predictions)
     conf_matrix = confusion_matrix(Y_test, predictions)
     class_report = classification_report(Y_test, predictions)
     print("Accuracy: ", accuracy)
     print("Confusion Matrix:\n", conf_matrix)
     print("Classification Report:\n", class_report)

```

Accuracy: 0.8666666666666667

Confusion Matrix:

```

[[204  11]
 [ 29  56]]

```

Classification Report:

	precision	recall	f1-score	support
0.0	0.88	0.95	0.91	215
1.0	0.84	0.66	0.74	85
accuracy			0.87	300
macro avg	0.86	0.80	0.82	300
weighted avg	0.86	0.87	0.86	300

```

[ ]: X = df_min_max.iloc[:, :-1].values
     Y = df_min_max.iloc[:, -1].values.reshape(-1,1)
     from sklearn.model_selection import train_test_split
     X_train_dt, X_test_dt, Y_train_dt, Y_test_dt = train_test_split(X, Y, \
     ↪test_size=.20, random_state=41)

```

```

[ ]: print(X_train_dt.shape)
     print(Y_train_dt.shape)

```

(1200, 10)

(1200, 1)

```
[ ]: classifier = DecisionTreeClassifier(min_samples_split=3, max_depth=6)
classifier.fit(X_train_dt, Y_train_dt)
classifier.print_tree()
```

```
X_9 <= 0.0 ? 0.14172068580101083
  left:X_7 <= 0.47 ? 0.0673748031138045
    left:X_3 <= 0.3333333333333333 ? 0.10481356777653078
      left:X_6 <= 0.79 ? 0.10587154972181712
        left:X_8 <= 0.73 ? 0.048765112141546496
          left:X_8 <= 0.0 ? 0.05124653739612173
            left:1.0
            right:0.0
          right:X_2 <= 0.3333333333333333 ? 0.336620644312952
            left:X_0 <= 0.16666666666666666 ?
0.19753086419753085
                                                    left:1.0
                                                    right:0.0
                                                    right:1.0
          right:X_2 <= 0.3333333333333333 ? 0.14814814814814814
            left:X_8 <= 0.59 ? 0.49382716049382713
              left:0.0
              right:1.0
            right:1.0
          right:X_2 <= 0.3333333333333333 ? 0.06311638043195417
            left:X_8 <= 0.58 ? 0.18127940320899222
              left:X_6 <= 0.7 ? 0.43102040816326526
                left:0.0
                right:1.0
              right:X_7 <= 0.0 ? 0.06816568047337264
                left:0.0
                right:X_5 <= 0.8581793054445147 ?
0.03679999999999998
                                                    left:1.0
                                                    right:0.0
          right:X_3 <= 0.4 ? 0.023772418058132447
            left:1.0
            right:X_7 <= 0.04 ? 0.01028466483011925
              left:X_4 <= 0.75 ? 0.31999999999999984
                left:1.0
                right:0.0
              right:1.0
            right:X_3 <= 0.3333333333333333 ? 0.022031077115505032
              left:X_6 <= 0.65 ? 0.07500000000000001
                left:X_8 <= 0.52 ? 0.2049011857707509
                  left:X_2 <= 0.3333333333333333 ? 0.39669421487603307
                    left:0.0
                    right:1.0
                  right:X_7 <= 0.5 ? 0.07578621756315554
```



```

left:0.0
right:X_2 <= 0.0 ? 0.014049586776859371
left:1.0
right:1.0

right:1.0
right:X_2 <= 0.0 ? 0.007891110806042358
left:X_5 <= 0.8192442904827852 ? 0.11932938856015768
left:X_7 <= 0.63 ? 0.041666666666666796
left:X_3 <= 0.6 ? 0.4444444444444444
left:0.0
right:1.0

right:1.0
right:0.0
right:1.0
right:X_2 <= 0.3333333333333333 ? 0.020658593641384004
left:X_6 <= 0.74 ? 0.003701210553126749
left:X_0 <= 0.0 ? 0.0009092584064248482
left:X_7 <= 0.66 ? 0.08357142857142866
left:0.0
right:X_7 <= 0.72 ? 0.48979591836734704
left:1.0
right:0.0
right:X_6 <= 0.21 ? 0.000690975656774756
left:X_6 <= 0.2 ? 0.012599726837922265
left:X_0 <= 0.9 ? 0.008895044629116652
left:0.0
right:0.0
right:X_0 <= 0.36666666666666664 ? 0.5
left:1.0
right:0.0

right:X_8 <= 0.77 ? 0.0007976258462526614
left:X_8 <= 0.32 ? 0.0008267812083850856
left:0.0
right:0.0
right:X_8 <= 0.8 ? 0.03481661515687966
left:1.0
right:0.0

right:X_8 <= 0.59 ? 0.04321229107081567
left:X_0 <= 0.1 ? 0.003245456378361794
left:X_3 <= 0.0 ? 0.0973183391003461
left:1.0
right:X_8 <= 0.42 ? 0.0234375
left:0.0
right:0.0

right:X_8 <= 0.03 ? 0.006156181480856789
left:X_0 <= 0.7 ? 0.375
left:0.0
right:1.0

```

```

right:0.0
right:X_7 <= 0.74 ? 0.14713482292141655
left:X_8 <= 0.62 ? 0.06838238203907196
left:1.0
right:X_9 <= 0.5 ? 0.016671173363618286
left:0.0
right:0.0
right:X_3 <= 0.2666666666666666 ? 0.26574394463667833
left:X_0 <= 0.0 ? 0.31999999999999984
left:1.0
right:0.0
right:1.0
right:X_8 <= 0.61 ? 0.08393651806734992
left:X_6 <= 0.71 ? 0.0439372843616922
left:X_5 <= 0.7890827161160316 ? 0.004056068742639862
left:X_4 <= 0.25 ? 0.003533208765075388
left:X_4 <= 0.0 ? 0.01387755102040833
left:0.0
right:0.0
right:0.0
right:X_0 <= 0.0 ? 0.07605263157894751
left:1.0
right:X_5 <= 0.7964441031675534 ?
0.08895044629116648
left:1.0
right:0.0
right:X_3 <= 0.4666666666666667 ? 0.24000000000000001
left:X_5 <= 0.0801045939366328 ? 0.08526315789473664
left:1.0
right:X_5 <= 0.9790540099529582 ?
0.09972299168975085
left:0.0
right:1.0
right:X_7 <= 0.29 ? 0.31999999999999984
left:0.0
right:1.0
right:X_3 <= 0.4 ? 0.19839106523237326
left:X_6 <= 0.71 ? 0.07772268135904487
left:X_5 <= 0.011299260249263768 ? 0.03679999999999998
left:0.0
right:0.0
right:X_7 <= 0.24 ? 0.5
left:0.0
right:1.0
right:X_7 <= 0.47 ? 0.11071983141082514
left:X_6 <= 0.52 ? 0.499054820415879
left:0.0
right:1.0

```

right:1.0

```
[ ]: from sklearn.metrics import accuracy_score, confusion_matrix, \
      ↪classification_report
prediction= classifier.predict(X_test_dt)
accuracy = accuracy_score(Y_test_dt, prediction)
conf_matrix = confusion_matrix(Y_test_dt, prediction)
class_report = classification_report(Y_test_dt,prediction)
print("Accuracy: ", accuracy)
print("Confusion Matrix:\n", conf_matrix)
print("Classification Report:\n", class_report)
```

Accuracy: 0.8966666666666666

Confusion Matrix:

[[192 16]

[15 77]]

Classification Report:

	precision	recall	f1-score	support
0.0	0.93	0.92	0.93	208
1.0	0.83	0.84	0.83	92
accuracy			0.90	300
macro avg	0.88	0.88	0.88	300
weighted avg	0.90	0.90	0.90	300

1 Naive Bayes

```
[ ]: NUM_FEATURES = 10
NUM_CLASSES = 2
NUM_POINTS = 1500
METRIC = "gaussian"
```

```
[ ]: def plotDists(y) :
      return sns.countplot(x=y)

def acc(true, pred) :
    assert len(true)==len(pred), "Truth and Pred Lengths not same"
    true = np.array(true)
    pred = np.array(pred).astype(np.int32)
    return np.sum(true==pred)/len(true)
```

```
[ ]: class NaiveBayes :
      def __init__(self, x, y, metric="gaussian", use_log=False) :
          self.x = np.array(x)
          self.y = np.array(y)
```

```

self.data = np.concatenate((self.x, self.y.reshape(-1, 1)), axis=1)
self.n_features = self.x.shape[1]
self.classes = set(y)
self.n_classes = len(self.classes)
assert metric in ["gaussian"], "Invalid metric"
self.metric = metric
self.use_log = use_log

def fit(self) :
    self.summary = {}
    for i in self.classes :
        d = self.data[self.data[:, -1] == i]
        mean = np.mean(d, axis=0)[:-1]
        std = np.std(d, axis=0)[:-1]
        l = len(d)
        self.summary[i] = {
            "mean" : mean,
            "std" : std,
            "len" : l
        }

def get_probability(self, inp, mean, std) :
    if self.metric == "gaussian" :
        exponent = np.exp(-(((inp - mean)**2)/(2*(std**2))))
        res = (1 / (np.sqrt(2 * np.pi) * std)) * exponent
        if self.use_log :
            return np.log(1 + res)
        return res

def predict(self, x) :
    assert self.summary, "Classifier not fit yet"
    results = []
    # Convert x to numerical type if it's a DataFrame
    if isinstance(x, pd.DataFrame):
        x = x.values.astype(np.float64)
    for inp in x :
        pred_class = -1
        pred_prob = 0
        for i in self.classes :
            probs = self.get_probability(inp, self.summary[i]["mean"], self.
↪summary[i]["std"])
            class_prob = np.prod(probs)
            if class_prob > pred_prob :
                pred_class = i
                pred_prob = class_prob
        results.append(pred_class)
    return results

```

```
[ ]: x = df_min_max.drop("HiringDecision",axis=1)
y = df_min_max["HiringDecision"]
df_min_max.head()
```


```
[ ]:      Age  Gender  EducationLevel  ExperienceYears  PreviousCompanies  \
0  0.200000    1.0      0.333333      0.000000      0.50
1  0.633333    1.0      1.000000      0.800000      0.50
2  0.933333    0.0      0.333333      0.200000      0.25
3  0.466667    1.0      0.333333      0.333333      0.25
4  0.333333    0.0      0.000000      0.400000      0.00
```

```
      DistanceFromCompany  InterviewScore  SkillScore  PersonalityScore  \
0           0.515450           0.48           0.78           0.91
1           0.497013           0.35           0.68           0.80
2           0.177927           0.20           0.67           0.13
3           0.107611           0.36           0.27           0.70
4           0.842135           0.23           0.52           0.85
```

```
      RecruitmentStrategy  HiringDecision
0           0.0           1.0
1           0.5           1.0
2           0.5           0.0
3           1.0           0.0
4           0.5           0.0
```


```
[ ]: nbg = NaiveBayes(x, y, METRIC)
nbg.fit()
acc(y, nbg.predict(x.values))
plotDists(nbg.predict(x))
```

```
[ ]: <Axes: ylabel='count'>
```



```
[ ]: nbgl = NaiveBayes(x, y, METRIC, use_log=True)
      nbgl.fit()
      acc(y, nbgl.predict(x))
      plotDists(nbgl.predict(x))
```

```
[ ]: <Axes: ylabel='count'>
```



```
[ ]: from sklearn.metrics import confusion_matrix

# Assuming 'y' contains the true labels and 'nbg.predict(x)' contains the
    ↪ predicted labels
y_pred = nbg.predict(x)
cm = confusion_matrix(y, y_pred)

print("Confusion Matrix:")
print(cm)

#Confusion Matrix
print("\nExplanation of the Confusion Matrix:")
print("True Positive (TP):", cm[1, 1], " - The model correctly predicted the
    ↪ positive class (HiringDecision = 1).")
print("True Negative (TN):", cm[0, 0], " - The model correctly predicted the
    ↪ negative class (HiringDecision = 0).")
print("False Positive (FP):", cm[0, 1], " - The model incorrectly predicted the
    ↪ positive class (HiringDecision = 1) when it was actually negative
    ↪ (HiringDecision = 0).")
```

```

print("False Negative (FN):", cm[1, 0], " - The model incorrectly predicted the
↳negative class (HiringDecision = 0) when it was actually positive
↳(HiringDecision = 1).")

# Additional analysis (optional)
accuracy = (cm[0, 0] + cm[1, 1]) / np.sum(cm)
precision = cm[1, 1] / (cm[1, 1] + cm[0, 1])
recall = cm[1, 1] / (cm[1, 1] + cm[1, 0])
f1_score = 2 * (precision * recall) / (precision + recall)

print("\nAccuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-Score:", f1_score)

```

Confusion Matrix:

```

[[839 196]
 [ 69 396]]

```

Explanation of the Confusion Matrix:

True Positive (TP): 396 - The model correctly predicted the positive class (HiringDecision = 1).

True Negative (TN): 839 - The model correctly predicted the negative class (HiringDecision = 0).

False Positive (FP): 196 - The model incorrectly predicted the positive class (HiringDecision = 1) when it was actually negative (HiringDecision = 0).

False Negative (FN): 69 - The model incorrectly predicted the negative class (HiringDecision = 0) when it was actually positive (HiringDecision = 1).

Accuracy: 0.8233333333333334

Precision: 0.668918918918919

Recall: 0.8516129032258064

F1-Score: 0.7492904446546831

```

[ ]: from sklearn.metrics import confusion_matrix

# Assuming 'y' contains the true labels and 'nbg.predict(x)' contains the
↳predicted labels
y_pred = nbg.predict(x)
cm = confusion_matrix(y, y_pred)

print("Confusion Matrix:")
print(cm)

#Confusion Matrix
print("\nExplanation of the Confusion Matrix:")

```



```

print("True Positive (TP):", cm[1, 1], " - The model correctly predicted the_
↳positive class (HiringDecision = 1).")
print("True Negative (TN):", cm[0, 0], " - The model correctly predicted the_
↳negative class (HiringDecision = 0).")
print("False Positive (FP):", cm[0, 1], " - The model incorrectly predicted the_
↳positive class (HiringDecision = 1) when it was actually negative_
↳(HiringDecision = 0).")
print("False Negative (FN):", cm[1, 0], " - The model incorrectly predicted the_
↳negative class (HiringDecision = 0) when it was actually positive_
↳(HiringDecision = 1).")

# Additional analysis (optional)
accuracy = (cm[0, 0] + cm[1, 1]) / np.sum(cm)
precision = cm[1, 1] / (cm[1, 1] + cm[0, 1])
recall = cm[1, 1] / (cm[1, 1] + cm[1, 0])
f1_score = 2 * (precision * recall) / (precision + recall)

print("\nAccuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-Score:", f1_score)

```

Confusion Matrix:

```

[[839 196]
 [ 69 396]]

```

Explanation of the Confusion Matrix:

True Positive (TP): 396 - The model correctly predicted the positive class (HiringDecision = 1).

True Negative (TN): 839 - The model correctly predicted the negative class (HiringDecision = 0).

False Positive (FP): 196 - The model incorrectly predicted the positive class (HiringDecision = 1) when it was actually negative (HiringDecision = 0).

False Negative (FN): 69 - The model incorrectly predicted the negative class (HiringDecision = 0) when it was actually positive (HiringDecision = 1).

Accuracy: 0.8233333333333334

Precision: 0.668918918918919

Recall: 0.8516129032258064

F1-Score: 0.7492904446546831

[11]: *# CONCLUSION:*

```

# Comparative analysis of Logistic Regression, Decision Tree, and Naive Bayes_
↳for hiring prediction reveals that Decision Tree consistently
# outperforms the other models.

```

Decision Tree achieved a notable accuracy of 89.6% compared to Logistic
→Regression's 86.7% and Naive Bayes' 82.3%.
These findings strongly suggest that Decision Tree is the optimal choice for
→predicting hiring outcomes in this specific dataset